BST 261: Data Science II Lecture 2 MLPs

Heather Mattie
Harvard T.H. Chan School of Public Health
Spring 2 2020

Recipe of the Day

Blueberry scones with lemon glaze



Deep learning glossaries

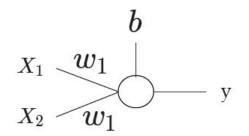
- 1. Google
- 2. WildML



Perceptrons

Perceptrons

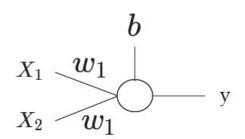
- Let's put this all together
- Our first network will be a single neuron that will learn a simple function



X1	X2	у
0	0	0
0	1	1
1	0	1
1	1	1

Perceptrons

O How do we make a prediction for each observation?



Assume the following values:

w1	w2	b
1	-1	-0.5

Observations

X1	X2	у
0	0	0
0	1	1
1	0	1
1	1	1

- \bigcirc For the first observation, $X_1 = 0, X_2 = 0, y = 0$
- First compute the weighted sum:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$h = 1 * 0 + -1 * 0 + (-0.5)$$

$$h = -0.5$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Observations

X1	X2	У
0	0	0
0	1	1
1	0	1
1	1	1

- \bigcirc For the first observation, $X_1 = 0, X_2 = 0, y = 0$
- First compute the weighted sum:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$h = 1 * 0 + -1 * 0 + (-0.5)$$

$$h = -0.5$$



$$p = \frac{1}{1 + \exp(-h)}$$

$$p = \frac{1}{1 + \exp(-0.5)}$$

$$p = 0.38$$

Assume the following values:

w1	w2	b		
1	-1	-0.5		

- \bigcirc For the first observation, $X_1 = 0, X_2 = 0, y = 0$
- First compute the weighted sum:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$h = 1 * 0 + -1 * 0 + (-0.5)$$

$$h = -0.5$$



$$p = \frac{1}{1 + \exp(-h)}$$

$$p = \frac{1}{1 + \exp(-0.5)}$$

$$p = 0.38$$



Assume the following values:

w1	w2	b
1	-1	-0.5

Round to get prediction:

$$\hat{y} = \text{round}(p)$$

$$\hat{y} = 0$$

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$\hat{y} = \text{round}(p)$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Complete the table:

					0
X1	X2	у	h	р	\hat{y}
0	0	0	-0.5	0.38	0
0	1	1			
1	0	1			
1	1	1			

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$\hat{y} = \text{round}(p)$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Complete the table:

					0
X1	X2	у	h	р	\hat{y}
0	0	0	-0.5	0.38	0
0	1	1	-1.5	0.18	0
1	0	1	0.5	0.38	0
1	1	1	-0.5	0.38	0

Performance

- Our network isn't so great
- O How do we make it better?
- What does better mean?
 - Need to define a measure of performance
 - There are many ways
- O Let's begin with squared error: $(y-p)^2$
- \bigcirc We need to find values for w_1, w_2, b that make this error as small as possible.
- We need to **learn** values for w_1, w_2, b such that the difference between the predicted and actual values is as small as possible.

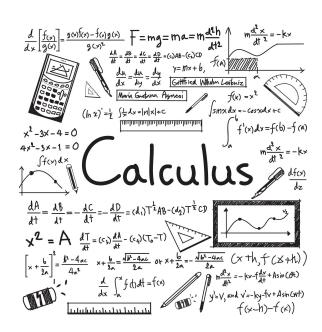
Learning From Data

How do we find the best values for w_1, w_2, b ?

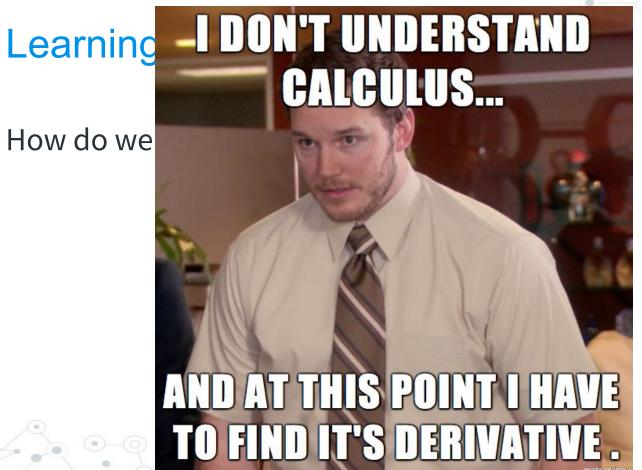


Learning From Data

How do we find the best values for w_1, w_2, b ?



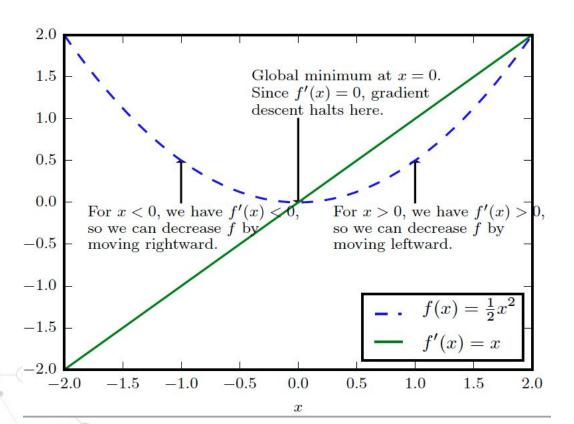
How do we



Learning From Data

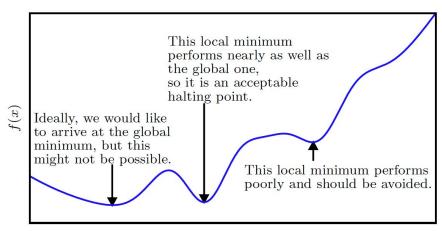
- Recall that the derivative of a function tells you how it is changing at any given location.
 - If the derivative is positive, it means it's going up
 - If the derivative is negative, it means it's going down
- Strategy:
 - \circ Start with initial values for w_1, w_2, b
 - \circ Take partial derivatives of the loss function with respect to w_1, w_2, b
 - Subtract the derivative (also called the gradient) from each
 - This is known as gradient descent

Gradient-Based Optimization



Gradient-Based Optimization

- \bigcirc A point that obtains the absolute lowest value of f(x) is a global minimum
- There may be one global minimum or multiple global minima
- It is also possible for there to be local minima that are not globally optimal
- It is common in many settings to settle for a value f that is very low but not necessarily minimal



Gradient-Based Optimization

- To minimize f, we would like to find the direction in which f decreases the fastest
- It can be shown that the gradient points directly uphill and the negative gradient directly downhill
- We can therefore decrease f by moving in the direction of the negative gradient
- \odot For example, for a weight w_i

$$w_i^{
m new} = w_i^{
m old} - \eta g$$

where $\,\eta\,$ is the **learning rate** (how fast you want to move down the gradient), and g is the gradient

Learning representations by back-propagating errors

Pittsburgh, Philadelphia 15213, USA

David E. Rumelhart*, Geoffrey E. Hinton† & Ronald J. Williams*

* Institute for Cognitive Science, C-015, University of California, San Diego, La Jolla, California 92093, USA † Department of Computer Science, Carnegie-Mellon University,

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure.



Our perceptron performs the following computations:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

We want to minimize this quantity:

$$l = (y - p)^2$$

We'll compute the gradients for each parameter by "backpropagating"
 errors through each component of the network

For w_1 we need to compute

$$rac{\partial l}{\partial w_1}$$

To get there we will use the chain rule

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$l = (y - p)^2$$

This is "backprop"

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} =$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$l = (y - p)^2$$

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2*(p-y)$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$l = (y - p)^2$$

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2*(p-y)$$

$$\frac{\partial p}{\partial h} =$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$l = (y - p)^2$$

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2*(p-y)$$

$$\left| \frac{\partial p}{\partial h} \right| = p * (1-p)$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$l = (y - p)^2$$

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2*(p-y)$$

$$\left| \frac{\partial p}{\partial h} \right| = p * (1-p)$$

$$\left| rac{\partial h}{\partial w_1}
ight| =$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$l = (y - p)^2$$

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2*(p-y)$$

$$\left| \frac{\partial p}{\partial h} \right| = p * (1 - p)$$

$$\frac{\partial h}{\partial w_1} = X_1$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$l = (y - p)^2$$

Putting it all together:

$$\frac{\partial l}{\partial w_1} = 2 * (p - y) * p * (1 - p) * X_1$$

Gradient Descent with Backprop

For some number of iterations:

- 1. Compute the gradient for w_1
- 2. Update $w_i^{\text{new}} = w_i^{\text{old}} \eta g$
- 3. Repeat until "convergence"

Do this for each weight and bias term.

Multilayer Perceptrons

Perceptron → MLP

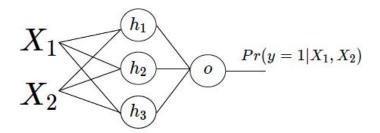
We can turn our perceptron model into a multilayer perceptron

- Instead of just one linear combination, we are going to take several, each with a different set of weights
- Each linear combination will be followed by a nonlinear activation
- Each of these nonlinear features will be fed into the logistic regression classifier (binary classifier)
- All of the weights are learned end-to-end via SGD

MLPs learn a set of nonlinear features directly from data - "feature engineering" is the hallmark of deep learning approaches

Multilayer Perceptrons (MLPs)

Suppose we have the following MLP with 1 hidden layer that has 3 hidden units:



Each neuron in the hidden layer is going to do exactly the same thing as before.

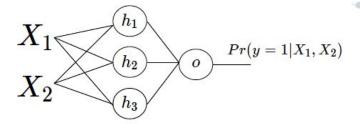
Multilayer Perceptrons (MLPs)

Computations:

$$h_{j} = \phi(w_{1j} * X_{1} + w_{2j} * X_{2} + b_{j})$$

$$o = b_{o} + \sum_{j=1}^{3} w_{oj} * h_{j}$$

$$p = \frac{1}{1 + \exp(-o)}$$



*If we use a sigmoid activation function

Output layer weight derivatives

$$\frac{\partial l}{\partial w_{oj}} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial o} * \frac{\partial o}{\partial w_{oj}}$$
$$= (p - y) * p * (1 - p) * h_j$$

Hidden layer weight derivatives

$$\frac{\partial l}{\partial w_{1j}} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial o} * \frac{\partial o}{\partial h} * \frac{\partial h}{\partial w_{1j}}$$
$$= (p - y) * p * (1 - p) * h_j * (1 - h_j) * X_1$$

Matrix Notation

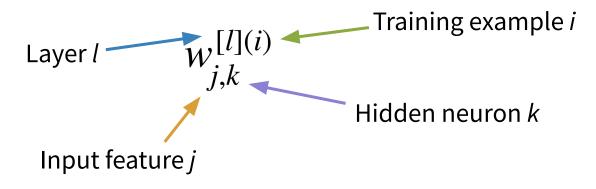
Sum notation starts to get unwieldy quickly. We can use matrix notation to represent each calculation in a more concise way.

$$X_1$$
 X_2
 h_2
 h_3
 O
 $Pr(y=1|X_1, X_2)$
 $X = \begin{bmatrix} X_1 \\ X_2 \end{bmatrix}$
 $W = \begin{bmatrix} w_{1,1} & w_{2,1} & w_{3,1} \\ w_{1,2} & w_{2,2} & w_{3,2} \end{bmatrix}$
 $B = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$
 $Z = W^T X + B$
 $H = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix} = \phi(Z)$

Notation

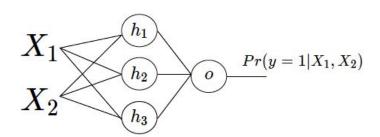
As the number of layers grows, the number of matrices grows and we have to add a superscript to denote the layer. We also have to add a superscript to denote which training example we are referencing.

Example notation for 1 weight in 1 hidden layer for 1 training example:



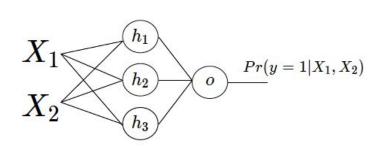
MLP Terminology

Forward pass = computing probability from input



MLP Terminology

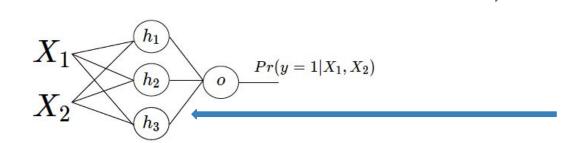
Forward pass = computing probability from input



Backward pass = computing derivatives from the output

MLP Terminology

Forward pass = computing probability from input

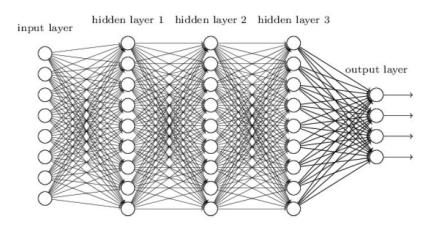


Hidden layers are also called "dense" layers or "fully connected" layers

Backward pass = computing derivatives from the output

MLPs

Increasing the number of layers increases the flexibility of the model - but run the risk of overfitting



Conclusions

 Backprop, perceptrons, and MLPs are the building blocks of neural nets

- You'll get a chance to demonstrate your mastery in Problem Set 1
- We will use these concepts for the rest of the semester

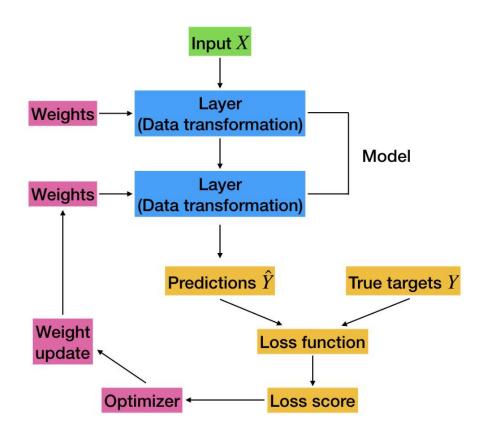
Coding Neural Nets

Keras and Tensorflow



- Keras is a model-level library that provides high-level building blocks for developing deep learning models
- It doesn't handle low-level operations like matrix and tensor (n-dimensional matrix) multiplication and differentiation
 - It uses TensorFlow or Theano or CNTK (Microsoft Cognitive Toolkit) backends for this
 - We will be using TensorFlow
 - It is the most widely adopted, scalable and production ready
- Keras can run on both CPUs and GPUs
 - When running on CPUs, uses Eigen for tensor operations
 - When running on GPUs, uses the NVIDIA CUDA Deep Neural Network library (cuDNN)

Neural Network Workflow



Generic Feedforward Network

Elements needed:

- 1. Necessary libraries
- 2. Dataset split into training and test sets (validation as well if you have enough data)
- 3. **models.sequential():** defines a linear, or sequential architecture made up of a set of layers that will stack to create the network
- 4. layers.Dense(): specifies a fully connected layer
- 5. **model.compile(optimizer, loss, metrics):** specifies how to execute the training of the network
- 6. model.fit(train_data, train_labels, epochs, batch_size):
 fits the neural net using the training data, runs for a specified number of iterations
 (epochs) using batch_size number of training examples at a time

Generic Feedforward Network

```
1 # Load data
 2 (x train, y train), (x test, y test) = load data()
 4 # Define model
 5 # Start with linear stack of layers
 6 model = tf.keras.models.Sequential([
    # Layer 1 (Hidden layer, fully connected)
   tf.keras.layers.Dense(c, activation = 'activation function'),
    # Layer 2 (Output layer, fully connected)
    tf.keras.layers.Dense(d, activation = 'output activation function')
11 ])
12
13 # Define how to execute training
14 model.compile(optimizer = 'optimizing algorithm',
15
                 loss = 'loss function',
16
                metrics = ['performance metric'])
18 # Train the network
19 model.fit(x train, y train, epochs = e, batch size = b)
```

```
4 # Define model
                                        5 # Start with linear stack of layers
                                        6 model = tf.keras.models.Sequential([
                                            # Layer 1 (Hidden layer, fully connected)
                                           tf.keras.layers.Dense(c, activation = 'activation function'),
                                            # Layer 2 (Output layer, fully connected)
                                       10 tf.keras.layers.Dense(d, activation = 'output activation function')
                                       11 1)
                                       12
                                       13 # Define how to execute training
                                       14 model.compile(optimizer = 'optimizing algorithm',
                                                         loss = 'loss function',
                                       15
             Input X
                                                         metrics = ['performance metric'])
                                       16
                                       17
              Laver
Weights
                                       18 # Train the network
         (Data transformation)
                                       19 model.fit(x train, y train, epochs = e, batch size = b)
                            Model
              Laver
Weights
         (Data transformation)
           Predictions \hat{Y}
                           True targets Y
Weight
                   Loss function
update
      Optimizer -
                    Loss score
```

1 # Load data

3

2 (x train, y train), (x test, y test) = load data()

```
Input X
                     Layer
Weights
              Data transformation)
                                          Model
                     Laver
Weights
              (Data transformation)
                 Predictions \hat{Y}
                                        True targets Y
Weight
                             Loss function
update
         Optimizer -
                              Loss score
```

```
1 # Load data
 2 (x train, y train), (x test, y test) = load data()
 4 # Define model
 5 # Start with linear stack of layers
 6 model = tf.keras.models.Sequential([
    # Layer 1 (Hidden layer, fully connected)
    tf.keras.layers.Dense(c, activation = 'activation function'),
    # Layer 2 (Output layer, fully connected)
   tf.keras.layers.Dense(d, activation = 'output activation function')
11 1)
12
13 # Define how to execute training
14 model.compile(optimizer = 'optimizing algorithm',
                 loss = 'loss function',
15
                metrics = ['performance metric'])
16
17
18 # Train the network
19 model.fit(x train, y train, epochs = e, batch size = b)
```

```
Input X
                    Laver
Weights
             (Data transformation)
                                         Model
                    Laver
Weights
             (Data transformation)
                Predictions \hat{Y}
                                       True targets Y
Weight
                            Loss function
update
        Optimizer
                             Loss score
```

```
1 # Load data
 2 (x train, y train), (x test, y test) = load data()
 3
 4 # Define model
 5 # Start with linear stack of layers
 6 model = tf.keras.models.Sequential([
     # Layer 1 (Hidden layer, fully connected)
    tf.keras.layers.Dense(c, activation = 'activation function'),
    # Layer 2 (Output layer, fully connected)
10 tf.keras.layers.Dense(d, activation = 'output activation function')
11 1)
12
13 # Define how to execute training
14 model.compile(optimizer = 'optimizing algorithm',
                 loss = 'loss function',
15
16
                 metrics = ['performance metric'])
17
18 # Train the network
19 model.fit(x train, y train, epochs = e, batch size = b)
```

```
2 (x train, y train), (x test, y test) = load data()
                                        3
                                       4 # Define model
                                        5 # Start with linear stack of layers
                                       6 model = tf.keras.models.Sequential([
                                           # Layer 1 (Hidden layer, fully connected)
                                          tf.keras.layers.Dense(c, activation = 'activation function'),
                                           # Layer 2 (Output layer, fully connected)
                                      10 tf.keras.layers.Dense(d, activation = 'output activation function')
                                      11 1)
                                      12
                                      13 # Define how to execute training
                                      14 model.compile(optimizer = 'optimizing algorithm',
                                                        loss = 'loss function',
                                      15
             Input X
                                                        metrics = ['performance metric'])
                                      16
                                      17
              Laver
Weights
                                      18 # Train the network
         (Data transformation)
                                      19 model.fit(x train, y train, epochs = e, batch size = b)
                           Model
              Laver
Weights
         (Data transformation)
           Predictions \hat{Y}
                          True targets Y
```

1 # Load data

Weight

update

Optimizer -

Loss function

Loss score

Generic Feedforward Network

train_data: training examples (matrix of feature vectors; **X**_{train})

train_labels: training labels (y_{train})

test_data: test examples used to measure performance of network (X_{test})

test_labels: test set labels (y_{test})

Optimizing algorithms: rmsprop, sgd, adagrad, adam, etc.

Loss function options: mse, mae, categorical_crossentropy, etc.

Performance measure options: accuracy, mae, etc.

Here:

c = the number of hidden units (neurons) in a hidden layer

d = the number of units (neurons) in the output layer

e = the number of epochs (iterations) over entire training data set

b = the batch size (how many training examples to optimize at once)